

Machine Learning Basics

Carlos Valle

Departamento de Informática
Universidad Técnica Federico Santa María

cvalle@inf.utfsm.cl

September 14, 2018

Overview

- 1 Introduction
- 2 Learning from examples
- 3 Linear Regression
- 4 Linear Classification
- 5 Logistic Regression
- 6 Cross validation
- 7 Curse of dimensionality
- 8 Shrinkage of coefficients
- 9 Shrinkage of coefficients

What is Machine Learning?

Arthur Samuel

“Machine Learning is the field of study that gives computers the ability to learn without being explicitly programmed”

Example: Dog recognizer

- Dog or no dog?



Example: Dog recognizer (2)

- Dog or no dog?



Example: Dog recognizer (3)

- Dog or no dog?



Training versus programming

Algorithm 1 Dog or no dog pseudocode

```
1: if four_legs & whiskers then  
2:   return dog  
3: else  
4:   return no dog  
5: end if
```



Training versus programming

Dogs



Training versus programming




What is Machine Learning?

Tom Mitchell

“A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E ”

Example

- Spam Detection



Example (2)

- Spam Detection

Subject	Sender	Date	Size
some there's a million miles to go...	Karen Burr	10/24/05 9:52	4KB
Want Free movies? Here's the answer!	Janis Whitley	12/30/04 11:02	6KB
Fwd: How to Get Your Free Adult Movies...	Andres Jimenez	12/23/04 8:45	5KB
No, get entire vhsd vcrs...	Patricia Hagen	12/17/04 5:32	2KB
no waiting room in clinic...	Marion Sanders	12/12/04 7:41	2KB
1/2 off new weekend pill so flavkfish...	Russell Lynn	11/15/04 11:24	1KB
Don't Wait Another Moment! Get Cash M...	Seasonal Daily Savings	10/18/04 5:17	1KB
Don't Wait Another Moment! Get Cash M...	Seasonal Daily Savings	10/18/04 5:16	1KB
Don't Wait Another Moment! Get Cash M...	Seasonal Daily Savings	10/18/04 5:16	1KB
The next time that could happen...	For Information	10/10/04 4:41 PM	12KB
No Text message	Parking	10/10 AM	2KB
HOCKEY 1991	Paula-Gail	9/10/04 10:42 AM	6KB
53--need a loan? it takes 2 minutes to...	Clara W. Goodwin	9/15/04 10:12	1KB
Health Insurance	Ameri Net	9/15/04 10:00	1KB
With I have to do all the paperwork myself, No...	Kali Sted	9/15/04 8:27 PM	1KB
Get your Associate's, Bachelor's, Master's...	Degree Match	9/15/04 5:14 PM	4KB
Remember me, Brian BODDIPATI...	Yan	9/15/04 1:09 PM	2KB
95% Off The Retailing Price For Saffo...	Alyson Swanson	9/15/04 10:30	5KB
Canta Rice Land Sale	Catholicism Properties	9/15/04 8:48 AM	3KB
Feel Yourself!! It's 747-819714	Low (bhatt)	9/15/04 8:40 AM	2KB
Get the Health care you deserve	Health Care	9/15/04 1:02 AM	2KB
There's no place like home	castlesn Grio	9/14/04 9:52 PM	1KB
How much are you paying for your cab...	Kerry Parkam	9/14/04 4:49 PM	2KB
Re: what do you think about this advert...	Mr. Hutchins	9/14/04 12:47	2KB
Credit Card Machines	Wholesale	9/14/04 11:45	3KB
Credit Card Machines	Wholesale	9/14/04 11:45	3KB
Credit Card Machines	Wholesale	9/14/04 11:45	3KB
This is what you were waiting for	Berna Corra	9/14/04 11:12	1KB
Single Good Day, hopefully you can see it	Bernie Tompkins	9/14/04 10:02	2KB
Your products will sell fast on eBay!!	Free-eBay-Seminar	9/14/04 9:16 AM	2KB
Get a Free Sony VAIO Laptop!	Limited Promotion	9/14/04 8:40 AM	4KB
Forum entry	Parking	9/14/04 8:10 AM	2KB
Re: MBA-qualificationkeyword	Efren Hyde	9/14/04 7:35 AM	5KB

SPAM

Example (3)

- Spam Detection

Subject	Sender	Date	Size
you there's a million miles to go... and	Karen Dwyer	10/24/04 9:52 ...	4KB
Want Free money? Here's the answer!	Joan Whiting	12/28/04 11:02...	5KB
Fwd: How to Get Your Free Adult Movies...	Andre Jamison	12/23/04 8:45 ...	5KB
Re: go online whadit waiyo sof	Pamela Hayes	12/17/04 5:52 ...	2KB
no waiting room in clinic tynek	Marion Sanders	12/12/04 7:41 ...	2KB
1/2 off new weekend pitc so Hawkfish...	Shadell Lyons	11/15/04 11:24...	1KB
Don't Wait Another Moment! Get Cash N...	Sensational Daily Savings	10/18/04 5:17 ...	1KB
Don't Wait Another Moment! Get Cash N...	Sensational Daily Savings	10/18/04 5:16 ...	1KB
You never know what could happen	My Responsibility	10/10/04 4:41 PM	1KB
Re: Test message	Parking	10-15 AM	2KB
53-and-a-bead it takes 2 minutes to	Clara W. Goodwin	9/15/04 10:12 ...	1KB
Health Insurance	Ameri Net	9/15/04 10:08 ...	1KB
Get your Associate's, Bachelor's, Master's	Dugan Plach	9/15/04 5:14 PM	4KB
Remember me, Mike (202)709551 from Mike	Ken	9/15/04 1:08 PM	2KB
95% Off The Retailing Prices For Selfw...	Alyson Deason	9/15/04 10:30 ...	5KB
Costa Rica Land Sale	Catalitan Properties	9/15/04 8:48 AM	3KB
Feel Yourself Fit + Tuff-877w	Levi DeHart	9/15/04 8:40 AM	2KB
Get the Health care you deserve	Health Care	9/15/04 1:02 AM	2KB
There's no place like home	Costless Eric	9/14/04 9:52 PM	1KB
How much are you paying for your cab!	Kerry Parkan	9/14/04 4:49 PM	2KB
Re: what do you think about this advert...	Mr. Hutchins	9/14/04 12:47 ...	2KB
Credit Card Machines	Wholesale	9/14/04 11:45 ...	3KB
Credit Card Machines	Wholesale	9/14/04 11:45 ...	3KB
Credit Card Machines	Wholesale	9/14/04 11:45 ...	3KB
This is what you were waiting for	Sarae Corra	9/14/04 11:12 ...	1KB
dinge Good Day, hopefully you can use 1...	beris templates	9/14/04 10:02 ...	2KB
Your products will sell fast as eBay! I ...	Free eBAY-Seminar	9/14/04 9:16 AM	2KB
Get a Free New Xbox Laptop!	Limited Promotion	9/14/04 8:47 AM	4KB
Re: MBA-qualification	Parking	10-15 AM	2KB
Re: MBA-qualification	Efren Hyde	9/14/04 7:35 AM	5KB

NO SPAM

Example (4)

- Spam Detection



Introduction to learning from examples

- Suppose we have a dataset giving weight, gender and calorie consumption a day from 40 people.

Weight (Kg)	Gender (M/F)	calorie cons. (cal)
85	M	2075
58	F	1757
52	M	2783
55	F	3500
\vdots	\vdots	\vdots

- We would like to predict the calorie consumption per day of other people.
- In this case, weight and gender are called **input features**. Each **vector input** \mathbf{x}_m has these two features.
- A calorie consumption per day y_m is called **target**.
- A pair (\mathbf{x}_m, y_m) is called a train example. A **train set** is a group of M training examples $S_M = \{(\mathbf{x}_m, y_m)\}, m = 1 \dots M$.

Preliminary definitions

- Features can be either **numerical** or **categorical**.
- Let \mathcal{X} be the **feature space**, and \mathcal{Y} the **output space**.
- Our goal is to obtain a mapping $f : \mathcal{X} \rightarrow \mathcal{Y}$, commonly called the *hypothesis* or **learner**), $\mathcal{X} \subseteq \mathbb{R}^n$, $y \in Y \subseteq \mathbb{R}$.
- Let \mathcal{S} the space that spans the possible samples, drawn from an unknown distribution $P(\mathbf{x}, y)$.
- A *learning algorithm* is a map from the space of train sets to the hypothesis space \mathcal{H} of possible functional solutions

$$\begin{aligned} \mathcal{A} &: \mathcal{S} \rightarrow \mathcal{H} \\ S_M &\rightarrow \mathcal{A}(S_M) = f. \end{aligned} \tag{1}$$

Training
Set


- If the target is continuous, we have a **regression problem**
- If the target can take a finite number k of discrete values, we have a **classification problem**. In particular $k = 2$ the problem is called **binary classification**.

Learning process




- If the target is continuous, we have a **regression problem**
- If the target can take a finite number k of discrete values, we have a **classification problem**. In particular $k = 2$ the problem is called **binary classification**.

Learning process




- If the target is continuous, we have a **regression problem**
- If the target can take a finite number k of discrete values, we have a **classification problem**. In particular $k = 2$ the problem is called **binary classification**.

Learning process



- If the target is continuous, we have a **regression problem**
- If the target can take a finite number k of discrete values, we have a **classification problem**. In particular $k = 2$ the problem is called **binary classification**.

Learning process



- If the target is continuous, we have a **regression problem**
- If the target can take a finite number k of discrete values, we have a **classification problem**. In particular $k = 2$ the problem is called **binary classification**.

Learning process (2)

- A main challenge is to construct an automatic method or **algorithm** able to estimate future examples based on the observed phenomenon in the train set.
- This key property of an algorithm is known as the **generalization ability**.
- The algorithms that memorize the train samples but have poor predictive performance with unknown examples, this undesirable problem is well-known as **overfitting**.

Loss function

- The quality of the algorithm \mathcal{A} is measured by the **loss function** given by $\ell : \mathbb{R} \times \mathcal{Y} \rightarrow [0, \infty)$, which quantifies the accuracy of the observed response $f(\mathbf{x})$ with respect to the true or **desired response** y .
- This function does not penalize the exact predictions, i.e., $\ell(y, f(\mathbf{x})) = 0$ if and only if $y = f(\mathbf{x})$.
- ℓ is a non-negative function, hence, the hypothesis will never profit from additional good predictions.
- In regression settings we use the **quadratic loss** function $\ell(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$
- While in classification we have the **misclassification loss** function

$$\ell(f(\mathbf{x}), y) = \begin{cases} 0 & \text{if } f(\mathbf{x}) = y \\ 1 & \text{if } f(\mathbf{x}) \neq y \end{cases}.$$

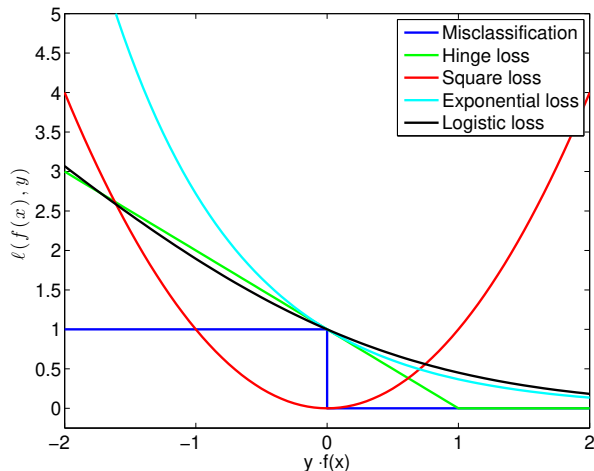
Loss function (2)

- However, In binary classification problem (where $y \in \{-1, 1\}$), the **margin** $yf(\mathbf{x})$ is introduced as a quality measure.
- This quality amount leads to several loss functions such as the *hinge loss*

$$\ell(f(\mathbf{x}), y) = \max(1 - yf(\mathbf{x}), 0) = |1 - yf(\mathbf{x})|_+.$$

- The *logistic loss* $\ell(f(\mathbf{x}), y) = \log_2 (1 + e^{-yf(\mathbf{x})})$
- The *exponential loss* $\ell(f(\mathbf{x}), y) = e^{-yf(\mathbf{x})}$.
- Note that the square loss can be arranged as $\ell(f(\mathbf{x}), y) = (y - f(\mathbf{x}))^2 = (1 - yf(\mathbf{x}))^2$ taking into account that $y^2 = 1$.
- And the misclassification loss can be written as $\ell(f(\mathbf{x}), y) = I(yf(\mathbf{x}) < 0)$, where $I(\cdot)$ is the **indicator function**.

Loss function (3)



- Logistic Loss and exponential loss can be viewed as a continuous approximation of the misclassification function

Motivation: Linear Regression

- We want to predict the total daily travel time of a trucker considering both the distance traveled and the number of deliveries made as input features:

Distance travel time (km)	Number of deliveries	Travel time (hr)
50	4	8.4
75	10	12.3
34	3	6.5
62	5	10.0
\vdots	\vdots	\vdots

- Here the inputs \mathbf{x} are bi-dimensional vectors. We let $\mathbf{x}_m^{(i)}$ denote the feature i of the m -th example.

- We could approximate the total daily travel time y as a linear function of the distance traveled and the number of deliveries made:

$$f(x) = \beta_0 + \beta_1 x^{(1)} + \beta_2 x^{(2)},$$

- where $\beta_i, i = 0, 1, 2$ are the parameters of the linear model.
- Thus, the space of linear functions mapping from \mathcal{X} to \mathcal{Y} is **parametric**.
- We define $x^{(0)} = 1$ to define the linear model in matrix form:

$$f(x) = \sum_{i=0}^I \beta_i x^{(i)} = \beta^T \mathbf{x}, \quad (2)$$

- where I is the number of features.

How do we select the β 's?

- We can get the parameters of the model by minimizing the quadratic loss function over the train set:

$$J(\beta) = \frac{1}{2} \sum_{m=1}^M (f(\mathbf{x}_m) - y_m)^2 \quad (3)$$

- where $\beta = (\beta_0, \beta_1, \dots, \beta_I)^T$.
- We need to choose β which minimizes $J(\beta)$.

- This problem can be expressed in matrix form:

$$J(\beta) = \frac{1}{2}(Y - X\beta)^T(Y - X\beta),$$

- where X is an $N \times (I + 1)$ matrix.

From matrix calculus

- If x is a column vector:
- $\frac{\partial u^T v}{\partial x} = \frac{\partial u}{\partial x} \cdot v + \frac{\partial v}{\partial x} \cdot u$
- $\frac{\partial Ax}{\partial x} = A^T$

- Then,

$$\begin{aligned}\frac{\partial J(\beta)}{\partial \beta} &= -X^T(Y - X\beta) \\ \frac{\partial^2 J(\beta)}{\partial \beta \partial \beta^T} &= X^T X\end{aligned}$$

- Equalizing the first derivative to zero we get the normal equations:

$$X^T(Y - X\beta) = 0 \implies \hat{\beta} = (X^T X)^{-1} X^T Y$$

Bias of the LMS Algorithm

- Assuming that the linear model is correct, i.e., $Y = X\beta + \epsilon$ for some unknown β . Furthermore, we assume that $E[\epsilon] = 0$ and $\text{Cov}(\epsilon) = E[\epsilon\epsilon^T] = \sigma^2 I$ (uncorrelated noise). From the least squares solution $\beta_{ls} = (X^T X)^{-1} X^T Y$. Thus

$$E[\beta_{ls}] = E\left[(X^T X)^{-1} X^T Y\right] \quad (4)$$

$$= (X^T X)^{-1} X^T E[Y] = (X^T X)^{-1} X^T X\beta = \beta, \quad (5)$$

- That is, β_{ls} is an unbiased estimator of β .
- Furthermore

$$\begin{aligned} \beta_{ls} - E[\beta_{ls}] &= (X^T X)^{-1} X^T Y - \beta \\ &= (X^T X)^{-1} X^T Y - (X^T X)^{-1} X^T X\beta \\ &= (X^T X)^{-1} X^T (Y - X\beta) \\ &= (X^T X)^{-1} X^T \epsilon \end{aligned} \quad (6)$$

Variance of the LMS Algorithm

- Thus, from the assumption $E[\epsilon\epsilon^T] = \sigma^2 I$ we obtain

$$\begin{aligned}\text{Cov}(\beta_{ls}) &= E \left[(\beta_{ls} - E[\beta_{ls}]) (\hat{\beta}^{ls} - E[\beta_{ls}])^T \right] \\ &= E \left[(X^T X)^{-1} X^T \epsilon \epsilon^T X (X^T X)^{-1} \right] \\ &= (X^T X)^{-1} X^T E[\epsilon \epsilon^T] X (X^T X)^{-1} \\ &= \sigma^2 (X^T X)^{-1} X^T X (X^T X)^{-1} \\ &= \sigma^2 (X^T X)^{-1}\end{aligned}\tag{7}$$

- Typically the variance σ^2 is estimated as

$$\hat{\sigma}^2 = \frac{1}{N - I - 1} \sum_{m=1}^M (y_m - \hat{y}_m)^2.$$

Gradient descent algorithm

- Let consider the **gradient descent** algorithm which start with some initial β and repeatedly performs:

$$\beta_i^{p+1} = \beta_i^p - \alpha \frac{\partial}{\partial \beta_i} J(\beta)$$

- Let's calculate the derivative of the loss function for a single train example (\mathbf{x}_m, y_m) :

$$\begin{aligned} \frac{\partial}{\partial \beta_i} J(\beta) &= \frac{\partial}{\partial \beta_i} \frac{1}{2} (f(\mathbf{x}_m) - y_m)^2 \\ &= (f(\mathbf{x}_m) - y_m) \cdot \frac{\partial}{\partial \beta_i} \left(\sum_{i=0}^I \beta_i x^{(i)} \right) \\ &= (f(\mathbf{x}_m) - y_m) \mathbf{x}_m^{(i)} \end{aligned}$$

Gradient descent algorithm (2)

- Then, for a single example:

$$\beta_i^{p+1} = \beta_i^p - \alpha(f(\mathbf{x}_m) - y_m) \mathbf{x}_m^{(i)}$$

- This rule is called Widrow-Hoff learning rule.
- Note that the amount of the update is proportional to the error:
 $(f(\mathbf{x}_m) - y_m)$

Batch gradient descent algorithm

- We can apply the learning rule above for the train set:

Algorithm 2 Batch gradient descent algorithm

```
1: repeat  
2:    $\beta_i^{p+1} = \beta_i^p - \alpha \sum_{m=1}^M (f(\mathbf{x}_m) - y_m) \mathbf{x}_m^{(i)}$  (for every  $i$ )  
3: until Convergence
```

- In general, gradient descent can reach a **local minimum**.
- However, J is a convex function. Thus, the optimization problem has only one **global optimum**.

Stochastic gradient descent algorithm

- We can modify the learning rule above for each example:



Algorithm 3 Stochastic gradient descent algorithm

```
1: repeat  
2:   for  $m := 1$  to  $M$  do  
3:      $\beta_i^{p+1} := \beta_i^p - \alpha(f(\mathbf{x}_m) - y_m) \mathbf{x}_m^{(i)}$  (for every  $i$ )  
4:      $\beta_i^p := \beta_i^{p+1}$   
5:   end for  
6: until Convergence
```

- This method is called **stochastic** or **online** gradient descent.
- Usually, this technique converge faster than batch gradient descent.
- However, using a fixed value for α it may never converge to the minimum of $J(\beta)$, oscillating around it.
- To avoid this behavior, it is recommended to slowly decrease α to zero along the iterations.

Why don't use a linear regression algorithm?

- Linear regression models might **mask** some classes.



- It measures the relationship between the class and the input vector by estimating probabilities using a logistic function:

-

$$f_{\beta}(\mathbf{x}) = g(\beta^T \mathbf{x}) = \frac{1}{1 + e^{-\beta^T \mathbf{x}}},$$


- where

$$g(z) = \frac{1}{1 + e^{-z}} = \frac{e^z}{1 + e^z}$$

is the **logistic** or **sigmoid** function

- As before, $\beta^T \mathbf{x} = \beta_0 + \sum_{i=1}^I \beta_i x(i)$
- However, it is worth mentioning that $f_{\beta}(\mathbf{x})$ is not linear.

Logistic function



- When $z = 0 = \beta^T x$, $g(z)$ is on the decision threshold 0.5.
- When $z \rightarrow \infty$, $g(z) \rightarrow 1$.
- When $z \rightarrow -\infty$, $g(z) \rightarrow 0$.

Logistic function (2)

- The derivative of $g(z)$ is
-

$$\begin{aligned}g'(z) &= \frac{d}{dz} \frac{1}{1 + e^{-z}} \\&= \frac{1}{(1 + e^{-z})^2} (e^{-z}) \\&= \frac{1}{(1 + e^{-z})} \cdot \left(1 - \frac{1}{(1 + e^{-z})}\right) \\&= g(z)(1 - g(z))\end{aligned}$$

How to obtain β ?

- Let us assume that

$$P(y = 1|\mathbf{x}; \beta) = f_{\beta}(\mathbf{x})$$

$$P(y = 0|\mathbf{x}; \beta) = 1 - f_{\beta}(\mathbf{x})$$

- Thus $p(y|\mathbf{x}; \beta) = (f_{\beta}(\mathbf{x}))^y(1 - f_{\beta}(\mathbf{x}))^{1-y}$

How to obtain β ? (2)

- Assuming that the M training example are independent, we can express the likelihood function as:

$$\begin{aligned} L(\beta) &= \prod_{m=1}^M p(y_m | \mathbf{x}_m; \beta) \\ &= \prod_{m=1}^M (f_{\beta}(\mathbf{x}_m))^{y_m} (1 - f_{\beta}(\mathbf{x}_m))^{1-y_m} \end{aligned}$$

- As we did before, we maximize the log likelihood function:

$$\begin{aligned} \ell(\beta) &= \sum_{m=1}^M y_m \log f_{\beta}(\mathbf{x}_m) + (1 - y_m) \log(1 - f_{\beta}(\mathbf{x}_m)) \\ &= \sum_{m=1}^M y_m \beta^T \mathbf{x}_m - \log(1 + e^{\beta^T \mathbf{x}_m}) \end{aligned} \tag{8}$$

Stochastic gradient ascent

- Now we can use gradient descent algorithm:

$$\beta^{p+1} = \beta^p + \alpha \nabla_{\beta} \ell(\beta)$$

- Let's obtain the derivative:

$$\begin{aligned} \frac{\partial}{\partial \beta_i} \ell(\beta) &= \left(y \frac{1}{f_{\beta}(\mathbf{x})} - (1-y) \frac{1}{1-f_{\beta}(\mathbf{x})} \right) \frac{\partial}{\partial \beta_i} f_{\beta}(\mathbf{x}) \\ &= \left(y \frac{1}{f_{\beta}(\mathbf{x})} - (1-y) \frac{1}{1-f_{\beta}(\mathbf{x})} \right) f_{\beta}(\mathbf{x})(1-f_{\beta}(\mathbf{x})) \frac{\partial}{\partial \beta_i} \beta^T \mathbf{x} \\ &= (y(1-f_{\beta}(\mathbf{x})) - (1-y)f_{\beta}(\mathbf{x})) x^{(i)} \\ &= (y - f_{\beta}(\mathbf{x})) x^{(i)} \end{aligned}$$

- Thus, the stochastic gradient ascent rule is

$$\beta_i^{p+1} = \beta_i^p + \alpha (y - f_{\beta}(\mathbf{x})) x^{(i)}$$

How to select a model in practice?

- How to automatically choose the parameters of the model?
- We will assume we have a finite set of models $\mathcal{M} = \{M_1, \dots, M_d\}$
- In case that the parameter(s) is (are) continuous we can discretize it (them).
- As we stated before, choosing the parameters which minimizes the train error does not guarantee generalization.

Hold-out cross validation

- One option is to do the following:

Algorithm 4 Hold-out cross validation

- 1: Randomly split S into S_{train} (For instance 75% of the data) and S_{cv} (the remaining 25%). Where S_{cv} is the hold-out cross validation set.
 - 2: Train each model M_i on S_{train} only to get the hypothesis f_i .
 - 3: Select and output the hypothesis f_i with the smallest error $\hat{\epsilon}_{SCV}(f_i)$ on the hold out cross validation set.
-

- Thus, for each f_i we obtain a better estimation of the generalization error.
- And we can select the hypothesis with the minimum estimated generalization error.
- Usually the size of the hold out cross validation set is set between 1/4 and 1/3.
- We are losing data that we are not using for training. And we estimate the generalization error only from **one** cross validation set.

K -fold cross validation

Algorithm 5 K -fold cross validation

- 1: Randomly split S into K disjoint subsets of M/K training examples each. Let's call these subsets S_1, \dots, S_K .
 - 2: **for** each model M_i **do**
 - 3: **for** $j = 1, \dots, K$ **do**
 - 4: $f_{ij} \leftarrow \mathcal{A}_i(S \setminus S_j)$ (train on all data except S_j)
 - 5: $\hat{\epsilon}_{S_j}(f_{ij}) = \frac{1}{|S_j|} \sum_{(\mathbf{x}_m, y_m) \in S_j} \ell(y_m, f_{ij}(\mathbf{x}_m))$ (Compute the error over the validation test)
 - 6: **end for**
 - 7: $\hat{\epsilon}_{M_i} = \frac{1}{K} \sum_{k=1}^K \hat{\epsilon}_{S_k}(f_{ik})$ (Calculate the average of the validation error over the k folds)
 - 8: **end for**
 - 9: Pick the model M_i with the lowest $\hat{\epsilon}_{M_i}$
 - 10: $f \leftarrow \mathcal{A}_i(S)$ (train a hypothesis with all training data according to the model M_i).
 - 11: **Output:** f
-

- \mathcal{A}_i is the learning algorithm according to the model M_i


K -fold cross validation (2)

- Usually $K = 5$ or $K = 10$.
- When $K = M$ this method is called **leave-one-out cross validation**.
- In classification problems, sometimes the proportion of examples of each class are unequal. In this case, we can adapt cross validation in such a way that each fold has the same proportions. Thus, all folds will be equally unbalanced.
- This method is called **stratified K -fold cross validation**.

Curse of dimensionality

- In real applications of machine learning, we have to deal with input spaces of high dimensionality comprising many inputs variables.
- Consider the problem of labeling a point according to the labels of its neighbors.
- To tackle this problem, we will divide the input space in cells, and we will label the point as the class having the largest number of training points in the same cell.
- This approach has a serious problem when x is a high-dimensional vector.
- If we divide a region of a space into regular cells, then the number of such cells grows exponentially with the dimensionality of the space.
- Therefore, with an exponentially large number of cells we would need an exponentially large amount of training examples in order to ensure that the cells are not empty.

Curse of dimensionality (2)



- In regression problems we would like to model the problem by considering the dependencies among the input features:

$$f(\mathbf{x}, \beta) = \beta_0 + \sum_{i=1}^I \beta_i \mathbf{x}^{(i)} + \sum_{i=1}^I \sum_{j=1}^I \beta_{ij} \mathbf{x}^{(i)} \mathbf{x}^{(j)} + \sum_{i=1}^I \sum_{j=1}^I \sum_{k=1}^I \beta_{ijk} \mathbf{x}^{(i)} \mathbf{x}^{(j)} \mathbf{x}^{(k)}$$

Curse of dimensionality (3)

- When I increases, the number of coefficients of β grows proportionally to I^3 .
- Note that we only considered up to 3 attributes. Thus, this method is not scalable.
- Now, consider a sphere of radius $r = 1$ in a space of I dimensions.
- The fraction of the volume of the sphere that lies between $r = 1 - \varepsilon$ and $r = 1$ is

$$\frac{V_I(1) - V_I(1 - \varepsilon)}{V_I(1)} = 1 - (1 - \varepsilon)^I,$$

- Here, $V_I(r) = K_I r^I$, where K_I is a constant that only depends on I .
- Thus, for high-dimensional spaces, the volume of the sphere is concentrated near to the surface.

Curse of dimensionality (4)

- Fortunately we can develop effective algorithm for high-dimensional data.
- Real Data is usually confined in a subset of the whole input space. So, we have lower effective dimensionality.
- Real data exhibits local smoothness properties. Hence, small changes in the input variables produces small changes in the target variables.
- We can exploit local interpolation-like techniques to predict the target of the new instances.

Bias-variance tradeoff

- A main challenge is to construct a learning algorithm able to **extrapolate**.
- That is, to estimate future examples based on the observed phenomenon in the train set. (**generalization error**)
- This key property of an algorithm is known as the **generalization ability**.
- On the other hand, we find the algorithms that memorize the training samples but have poor predictive performance with unknown examples, this undesirable problem is well-known as **overfitting**.
- If we compute a different model for each training sample that we have collected; the **bias** is the set of points that cannot be well predicted by the model, and the **variance** measures how different the predictions along the training samples are.

Bias-variance tradeoff (2)


- Mathematically,

$$\begin{aligned}E[(y - f(\mathbf{x}))^2] &= E[((y - E[f(\mathbf{x})]) + (E[f(\mathbf{x})] - f(\mathbf{x})))^2] \\&= E[(y - E[f(\mathbf{x})])^2] + E[2(y - E[f(\mathbf{x})])(E[f(\mathbf{x})] - f(\mathbf{x}))] \\&\quad + E[(E[f(\mathbf{x})] - f(\mathbf{x}))^2] \\&= E[(y - E[f(\mathbf{x})])^2] + E[(E[f(\mathbf{x})] - f(\mathbf{x}))^2] \\ \text{MSE} &= \text{bias}^2(f) + \text{var}(f)\end{aligned}$$

- From the machine learning point of view, this trade-off is strongly related with the complexity of the learner f .
- A learner with low complexity has high bias covering the training points, which can lead to [underfitting](#).

Bias-variance tradeoff (3)

- While, if the complexity of the learner is too high, the prediction tends to be closer (lower bias) to the training data and consequently generates overfitting.



Ridge regression

- It shrinks the coefficient values by introducing a regularization term,

$$\hat{\beta}^{ridge} = \operatorname{argmin}_{\beta} \left\{ \sum_{m=1}^M (y_m - \beta_0 - \sum_{i=1}^I x_{mi}\beta_i)^2 + \lambda \sum_{i=1}^I \beta_i^2 \right\} \quad (9)$$

- $\lambda \geq 0$ controls the regularization level
- Equivalently,

$$\begin{aligned} \hat{\beta}^{ridge} &= \operatorname{argmin}_{\beta} \sum_{m=1}^M \left(y_m - \beta_0 - \sum_{i=1}^I x_{mi}\beta_i \right)^2 \\ \text{s.a.} \quad &\sum_{i=1}^I \beta_i^2 \leq s \end{aligned}$$

- It is dependent of the scale of data. Thus, normalizing data is strictly necessary.
- β_0 is not bounded.

Reparametrizing

- We can transform data by using $x_{mi} \leftarrow (x_{mi} - \bar{x}_i)$
- Thus, we estimate $\beta_0 = \bar{y} = \sum_{m=1}^M y_m / M$
- Now X has I columns. Then, equation (9) can be expressed as

$$RSS(\lambda) = \frac{1}{2}((Y - X\beta)^T(Y - X\beta) + \lambda\beta^T\beta)$$

- The solution of this equation is

$$\hat{\beta}^{ridge} = (X^T X + \lambda I)^{-1} X^T Y$$

- Now the problem is non-singular.
- For orthonormal inputs we have, $\hat{\beta}_{ridge} = \gamma \hat{\beta}_{LS}, 0 \leq \gamma \leq 1$

Any questions?

